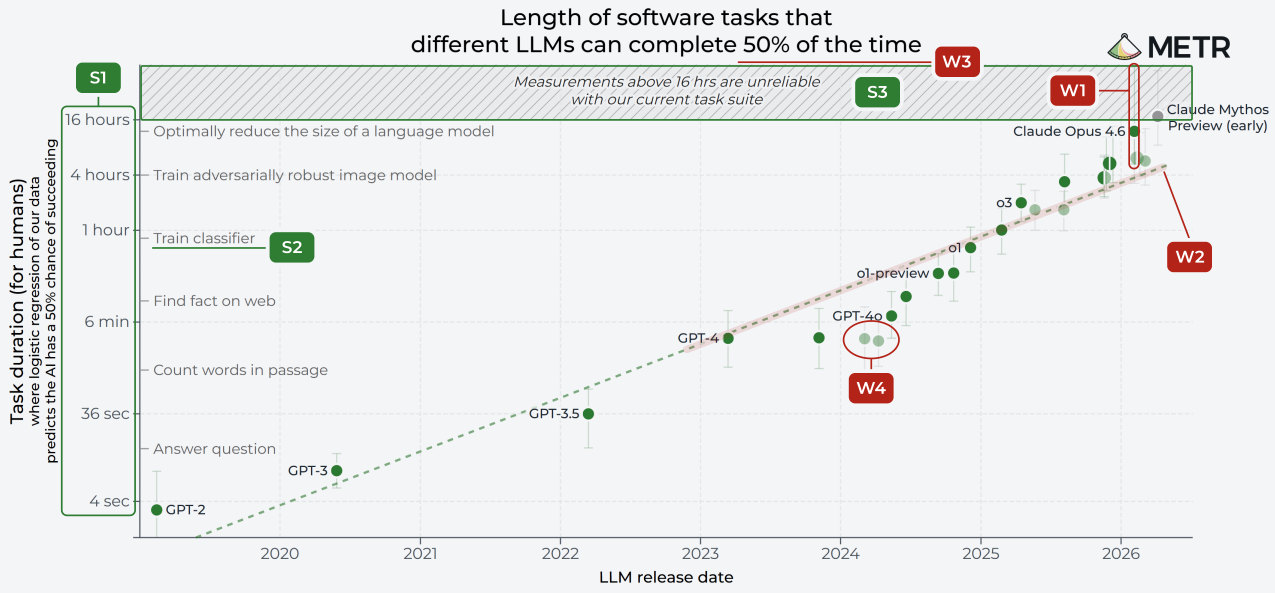
METR's Time Horizon Chart Is Easy to Read, and Easy to Overread

Valmik Nahata Varun Sinha Rishab Kolan Rishil Chintalapati
UC San Diego, DSC 106 Project 2.5 · {vnahata, v1sinha, rkolan, rchintalapati}@ucsd.edu

Original chart | Dataset | METR analysis code



METR, Time Horizon 1.1, 50% success, log scale. Our additions: S strength W weakness

Strengths

S1 Log scale makes doubling a straight line
Steady growth reads as a steady slope, through position on a common scale, the most accurate channel (L4). Ratios also fit Weber's law (L4).

S2 Example tasks give the axis meaning
"Train classifier" at about an hour turns minutes into kinds of work, a conceptual unit (data models, L2).

S3 It marks where its data stops
Hatching and the gray Mythos dot pop out preattentively (L4), so the chart claims nothing its tasks cannot measure (expressiveness, L2).

Weaknesses

W1 Uncertainty is drawn at 20% opacity
Opus 4.6 reads as 12 h, yet its 95% CI spans 5 to 61 h. Faint whiskers sink into the ground (L5). **Fix:** solid CI bars.

W2 One straight line hides a bend
Every frontier model since o3 sits above METR's line. Continuity (L5) reads it as one pace. **Fix:** add the recent fit.

W3 "50%" carries the headline
At 80% success Opus 4.6 drops from 12 h to 70 min, yet the threshold sits only in the title and a rotated label (L4). **Fix:** plot both.

W4 Light and split dots go unexplained
No key for light green, 10 of 26 models named, two models share one split dot (expressiveness, L2). **Fix:** key and jitter.

Marks and channels

Point	one AI model (26 in all)	x position	release date	Q interval
Vertical line	95% bootstrap CI	y position, log	50% horizon (human min)	Q ratio
Dashed line	METR's exponential fit	Whisker length	width of the CI	Q ratio
Hatched area	above 16 h	Luminance	frontier at release or not	N
Aa Text	model and task names	Gray hue, texture	beyond the 16 h ceiling	N

Our reanalysis

Since o3, every frontier model sits **above** METR's line. A fit from 2024 on doubles every **105 days** (METR: 188).

